

1) Starting the mysql by giving username as root and password as cloudera.

mysql -u root -pcloudera

```
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ mysql -u root -pcloudera
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 546
Server version: 5.1.73 Source distribution

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> █
```

show databases;

```
File Edit View Search Terminal Help
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| cm         |
| firehose   |
| hue        |
| metastore  |
| mysql      |
| nav        |
| navms      |
| oozie      |
| retail_db  |
| rman       |
| sentry     |
+-----+
12 rows in set (0.00 sec)

mysql> █
```

Use retail_db;

```
mysql> use retail_db;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> █
```

Show tables;

```
mysql> show tables;
+-----+
| Tables_in_retail_db |
+-----+
| categories           |
| customers            |
| departments          |
| order_items          |
| orders               |
| products             |
+-----+
6 rows in set (0.00 sec)

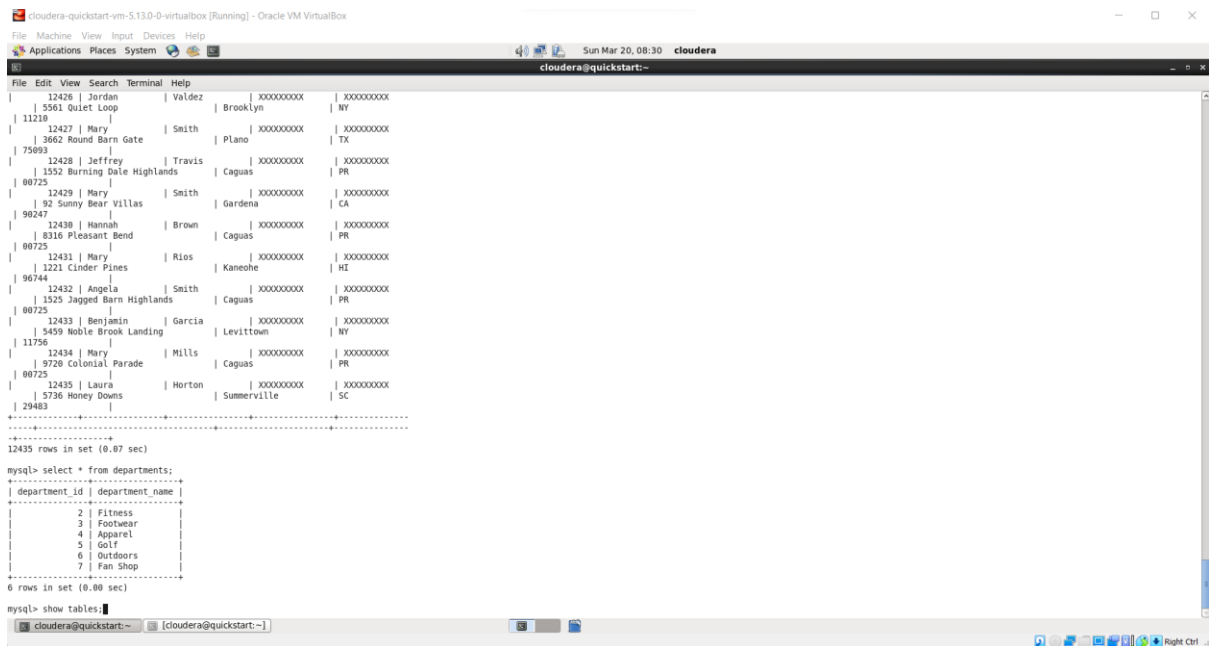
mysql>
```

Select * from customers;

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
Sun Mar 20, 08:28 cloudera
cloudera@quickstart:~$
File Edit View Search Terminal Help
12421 | Kyle | Small | XXXXXXXXXXXX | XXXXXXXXXXXX
2584 Round Square | Middletown | CT
06457 |
12422 | Mary | Blevins | XXXXXXXXXXXX | XXXXXXXXXXXX
6473 Bright Expressway | Caguas | PR
00725 |
12423 | Stephen | Smith | XXXXXXXXXXXX | XXXXXXXXXXXX
3445 Harvest Campus | Palmdale | CA
93550 |
12424 | Judy | Phillips | XXXXXXXXXXXX | XXXXXXXXXXXX
4534 Cinder Concession | San Diego | CA
92111 |
12425 | Mary | Smith | XXXXXXXXXXXX | XXXXXXXXXXXX
1858 Grand Forest Towers | Caguas | PR
00725 |
12426 | Jordan | Valdez | XXXXXXXXXXXX | XXXXXXXXXXXX
5561 Quiet Loop | Brooklyn | NY
11210 |
12427 | Mary | Smith | XXXXXXXXXXXX | XXXXXXXXXXXX
3662 Round Barn Gate | Plano | TX
75093 |
12428 | Jeffrey | Travis | XXXXXXXXXXXX | XXXXXXXXXXXX
1552 Burning Dale Highlands | Caguas | PR
00725 |
12429 | Mary | Smith | XXXXXXXXXXXX | XXXXXXXXXXXX
92 Sunny Bear Villas | Gardena | CA
90247 |
12430 | Hannah | Brown | XXXXXXXXXXXX | XXXXXXXXXXXX
8316 Pleasant Bend | Caguas | PR
00725 |
12431 | Mary | Rios | XXXXXXXXXXXX | XXXXXXXXXXXX
1221 Cinder Pines | Kaneohe | HI
96744 |
12432 | Angela | Smith | XXXXXXXXXXXX | XXXXXXXXXXXX
1525 Jagged Barn Highlands | Caguas | PR
00725 |
12433 | Benjamin | Garcia | XXXXXXXXXXXX | XXXXXXXXXXXX
5459 Noble Brook Landing | Levittown | NY
11756 |
12434 | Mary | Mills | XXXXXXXXXXXX | XXXXXXXXXXXX
9720 Colonial Parade | Caguas | PR
00725 |
12435 | Laura | Horton | XXXXXXXXXXXX | XXXXXXXXXXXX
5736 Honey Doves | Summerville | SC
29483 |
+-----+
12435 rows in set (0.07 sec)

cloudera@quickstart:~$
```

Select * from departments;



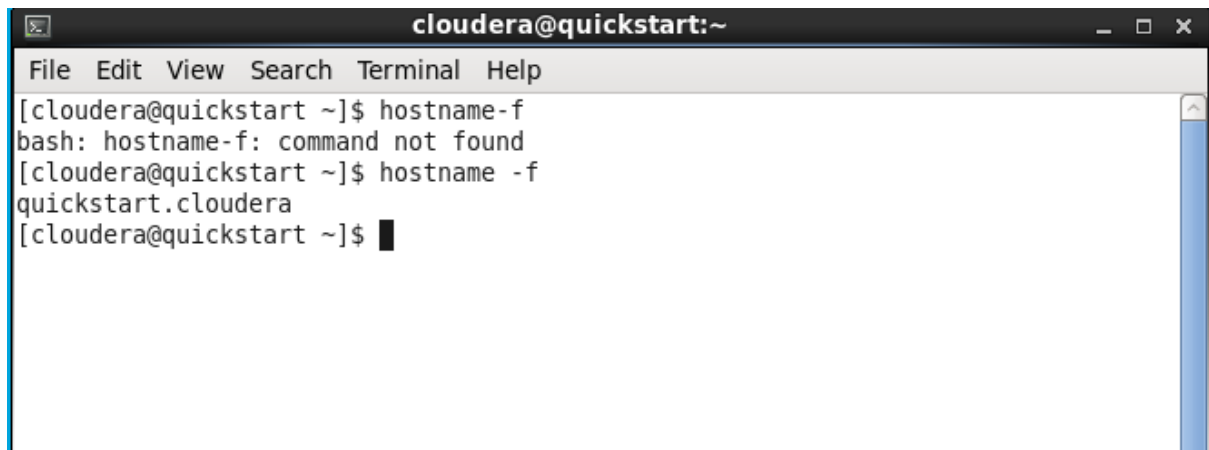
The screenshot shows a terminal window titled "cloudera@quickstart:~" with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal displays the output of a MySQL query: `mysql> select * from departments;`. The result is a table with 7 rows and 2 columns: `department_id` and `department_name`. The data is as follows:

department_id	department_name
2	Fitness
3	Footwear
4	Apparel
5	Golf
6	Outdoors
7	Fan Shop

Below the table, it says "6 rows in set (0.09 sec)". The terminal also shows the command `mysql> show tables;` which returns no results.

Open the new terminal for running command for sqoop.

Hostname-f



The screenshot shows a terminal window titled "cloudera@quickstart:~" with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal displays the output of the command `hostname -f`. The output is `quickstart.cloudera`.

`sqoop list-databases --connect jdbc:mysql://quickstart:3306/ --password cloudera --username root;`

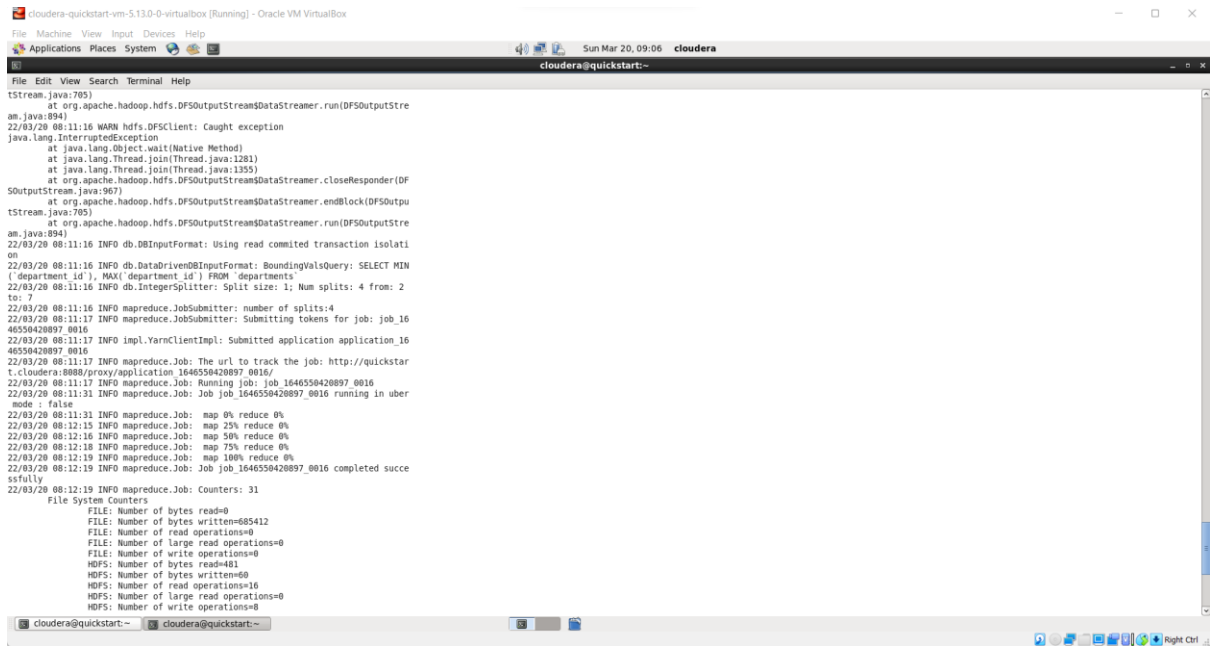
```
[cloudera@quickstart ~]$ sqoop list-tables --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root;
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
22/03/08 10:24:17 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/03/08 10:24:17 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
22/03/08 10:24:19 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
categories
customers
departments
order_items
orders
products
[cloudera@quickstart ~]$
```

sqoop list-tables --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table departments;

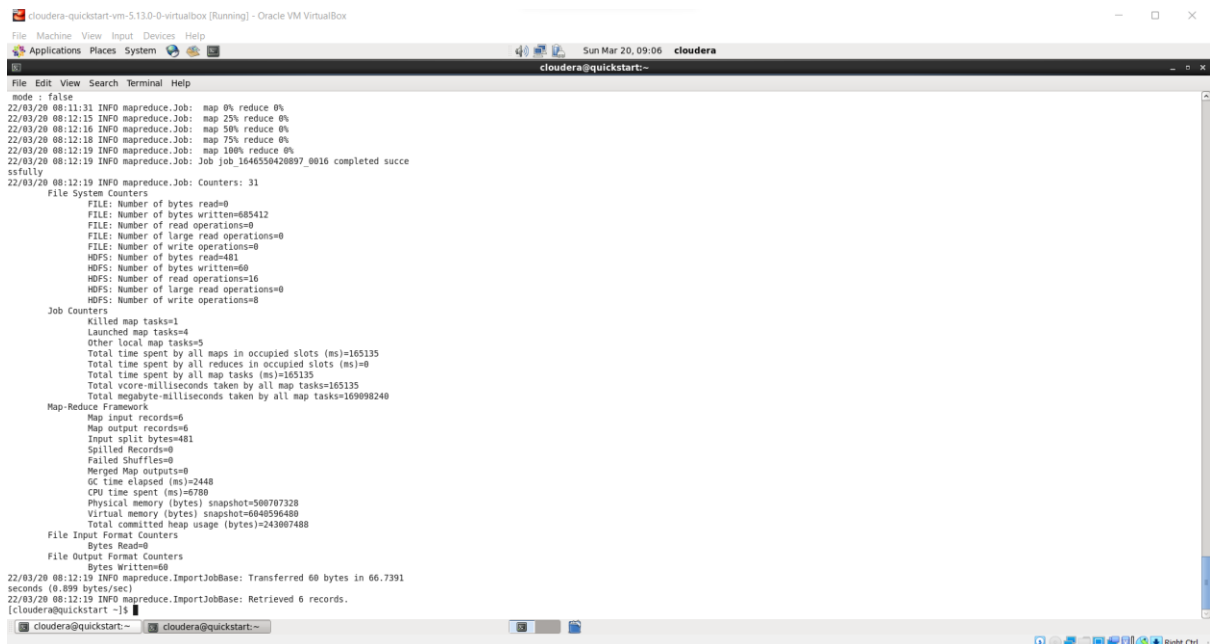


```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
Sun Mar 20, 08:54 cloudera
cloudera@quickstart:~$
File Edit View Search Terminal Help
oozie
retail_db
rman
sentry
[cloudera@quickstart ~]$ sqoop list-tables --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root;
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
22/03/20 07:55:33 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/03/20 07:55:33 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
22/03/20 07:55:34 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
categories
customers
departments
order_items
orders
products
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table departments;
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
22/03/20 07:58:39 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/03/20 07:58:39 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
22/03/20 07:58:40 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
22/03/20 07:58:40 INFO tool.CodeGenTool: Beginning code generation
22/03/20 07:58:41 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM departments AS t LIMIT 1
22/03/20 07:58:41 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM departments AS t LIMIT 1
22/03/20 07:58:41 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-cloudera/compile/az2b6727c73de13852a6820be248d920/departments.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
22/03/20 07:58:46 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/az2b6727c73de13852a6820be248d920/departments.jar
22/03/20 07:58:46 WARN manager.MySQLManager: It looks like you are importing from mysql.
22/03/20 07:58:46 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
22/03/20 07:58:46 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
22/03/20 07:58:46 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
22/03/20 07:58:46 INFO mapreduce.ImportJobBase: Beginning import of departments
```

```
cloudera-quickstart-vm-5.11.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$
File Edit View Search Terminal Help
at java.lang.Thread.join(Thread.java:1355)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFS
OutputStream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endLock(DFSOutpu
tStream.java:705)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at java.lang.Thread.join(Thread.java:1355)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFS
OutputStream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endLock(DFSOutput
Stream.java:705)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at java.lang.Thread.join(Thread.java:1355)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFS
OutputStream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeInternal(DFS
OutputStream.java:935)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:931)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at java.lang.Thread.join(Thread.java:1355)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFS
OutputStream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endLock(DFSOutput
Stream.java:705)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at java.lang.Thread.join(Thread.java:1355)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFS
OutputStream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endLock(DFSOutput
Stream.java:705)
```



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$
File Edit View Search Terminal Help
tstream.java:705)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:16 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DF
SOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutpu
tStream.java:705)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:16 INFO db.DBInputFormat: Using read committed transaction isolati
on
22/03/20 08:11:16 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN
('department id'), MAX('department id') FROM 'departments'
22/03/20 08:11:16 INFO db.IntegerSplitter: Split size: 1; Num splits: 4 from: 2
to: 7
22/03/20 08:11:16 INFO mapreduce.JobSubmitter: number of splits:4
22/03/20 08:11:17 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_16
46550420897_0016
22/03/20 08:11:17 INFO impl.YarnClientImpl: Submitted application application_16
46550420897_0016
22/03/20 08:11:17 INFO mapreduce.Job: The url to track the job: http://quickstar
t.cloudera:8080/proxy/application_1646550420897_0016/
22/03/20 08:11:17 INFO mapreduce.Job: Running job: job_1646550420897_0016
22/03/20 08:11:31 INFO mapreduce.Job: Job job_1646550420897_0016 running in uber
mode : false
22/03/20 08:11:31 INFO mapreduce.Job: map 0% reduce 0%
22/03/20 08:12:15 INFO mapreduce.Job: map 25% reduce 0%
22/03/20 08:12:16 INFO mapreduce.Job: map 50% reduce 0%
22/03/20 08:12:18 INFO mapreduce.Job: map 75% reduce 0%
22/03/20 08:12:19 INFO mapreduce.Job: map 100% reduce 0%
22/03/20 08:12:19 INFO mapreduce.Job: Job job_1646550420897_0016 completed succe
ssfully
22/03/20 08:12:19 INFO mapreduce.Job: Counters: 31
File System Counters
FILE: Number of bytes read=0
FILE: Number of bytes written=685412
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=481
HDFS: Number of bytes written=60
HDFS: Number of read operations=16
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
cloudera@quickstart:~$
```



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$
File Edit View Search Terminal Help
mode : false
22/03/20 08:11:31 INFO mapreduce.Job: map 0% reduce 0%
22/03/20 08:12:15 INFO mapreduce.Job: map 25% reduce 0%
22/03/20 08:12:16 INFO mapreduce.Job: map 50% reduce 0%
22/03/20 08:12:18 INFO mapreduce.Job: map 75% reduce 0%
22/03/20 08:12:19 INFO mapreduce.Job: map 100% reduce 0%
22/03/20 08:12:19 INFO mapreduce.Job: Job job_1646550420897_0016 completed succe
ssfully
22/03/20 08:12:19 INFO mapreduce.Job: Counters: 31
File System Counters
FILE: Number of bytes read=0
FILE: Number of bytes written=685412
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=481
HDFS: Number of bytes written=60
HDFS: Number of read operations=16
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
Job Counters
Killed map tasks=1
Launched map tasks=4
Other local map tasks=5
Total time spent by all maps in occupied slots (ms)=165135
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=165135
Total score-millisecods taken by all map tasks=165135
Total megabyte-millisecods taken by all map tasks=160990240
Map-Reduce Framework
Map input records=6
Map output records=6
Input split bytes=481
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=2448
CPU time spent (ms)=6780
Physical memory (bytes) snapshot=500707328
Virtual memory (bytes) snapshot=6849596480
Total committed heap usage (bytes)=243007488
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=60
22/03/20 08:12:19 INFO mapreduce.ImportJobBase: Transferred 60 bytes in 66.7391
seconds (0.899 bytes/sec)
22/03/20 08:12:19 INFO mapreduce.ImportJobBase: Retrieved 6 records.
cloudera@quickstart:~$
```

hadoop fs -ls

```
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=60
22/03/20 08:12:19 INFO mapreduce.ImportJobBase: Transferred 60 bytes in 66.7391
seconds (0.899 bytes/sec)
22/03/20 08:12:19 INFO mapreduce.ImportJobBase: Retrieved 6 records.
[cloudera@quickstart ~]$ hadoop fs -ls
Found 2 items
drwxr-xr-x - cloudera cloudera      0 2022-03-20 08:12 department1
drwxr-xr-x - cloudera cloudera      0 2022-03-20 07:59 departments
[cloudera@quickstart ~]$
```

hadoop fs -ls deparments;

```
[cloudera@quickstart ~]$ hadoop fs -ls
Found 1 items
drwxr-xr-x - cloudera cloudera      0 2022-03-20 07:59 departments
[cloudera@quickstart ~]$ hadoop fs -ls departments;
Found 5 items
-rw-r--r--  1 cloudera cloudera      0 2022-03-20 07:59 departments/_SUCCESS
;
-rw-r--r--  1 cloudera cloudera    21 2022-03-20 07:59 departments/part-m-
00000
-rw-r--r--  1 cloudera cloudera    10 2022-03-20 07:59 departments/part-m-
00001
-rw-r--r--  1 cloudera cloudera      7 2022-03-20 07:59 departments/part-m-
00002
-rw-r--r--  1 cloudera cloudera    22 2022-03-20 07:59 departments/part-m-
00003
```

hadoop fs -cat /user/cloudera/departments/part-m-00000

hadoop fs -cat /user/cloudera/departments/part-m-00001

hadoop fs -cat /user/cloudera/departments/part-m-00002

hadoop fs -cat /user/cloudera/departments/part-m-00003

```
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-00000
2,Fitness
3,Footwear
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-00001
4,Apparel
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-00002
5,Golf
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part-m-00006
cat: '/user/cloudera/departments/part-m-00006': No such file or directory
```

hadoop fs -cat /user/cloudera/departments/part*

```
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/departments/part*
2,Fitness
3,Footwear
4,Apparel
5,Golf
6,Outdoors
7,Fan Shop
```

sqoop import --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table departments --target-dir /user/cloudera/department1

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$ sqoop import --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table departments --target-dir /user/cloudera/department1
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set HADOOP_MAPRED_HOME to the root of your Accumulo installation.
22/03/20 08:11:03 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/03/20 08:11:03 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
22/03/20 08:11:04 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
22/03/20 08:11:04 INFO tool.CodeGenTool: Beginning code generation
22/03/20 08:11:05 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `departments` AS t LIMIT 1
22/03/20 08:11:05 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `departments` AS t LIMIT 1
22/03/20 08:11:05 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-cloudera/compile/Sed2741ae6e9b9d5fa8c2219d/departments.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
22/03/20 08:11:10 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/Sed2741ae6e9b9d5fa8c2219d/departments.jar
22/03/20 08:11:10 WARN manager.MySQLManager: It looks like you are importing from MySQL.
22/03/20 08:11:10 WARN manager.MySQLManager: This transfer can be faster! Use the --direct option to exercise a MySQL-specific fast path.
22/03/20 08:11:10 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
22/03/20 08:11:10 INFO mapreduce.ImportJobBase: Beginning import of departments
22/03/20 08:11:10 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.job.tracker.address
22/03/20 08:11:11 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
22/03/20 08:11:13 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
22/03/20 08:11:13 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/03/20 08:11:14 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1335)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DFSOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutputStream.java:705)
22/03/20 08:11:14 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1335)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DFSOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutputStream.java:705)
```



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$
File Edit View Search Terminal Help
an.java:831)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExceptio
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DF
OutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endLock(DFSOutpu
tStream.java:785)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExceptio
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DF
OutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endLock(DFSOutpu
tStream.java:785)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExceptio
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DF
OutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endLock(DFSOutpu
tStream.java:785)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExceptio
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DF
OutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endLock(DFSOutpu
tStream.java:785)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:16 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExceptio
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
```

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$
File Edit View Search Terminal Help
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutpu
tStream.java:705)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:15 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at java.lang.Thread.join(Thread.java:1355)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DF
SOutputStream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutpu
tStream.java:705)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:16 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at java.lang.Thread.join(Thread.java:1355)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DF
SOutputStream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutpu
tStream.java:705)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStre
am.java:894)
22/03/20 08:11:16 INFO db.DBInputFormat: Using read committed transaction isolati
on
22/03/20 08:11:16 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN
(department id), MAX(department id) FROM departments
22/03/20 08:11:16 INFO db.IntegerSplitter: Split size: 1; Num splits: 4 from: 2
to: 7
22/03/20 08:11:16 INFO mapreduce.JobSubmitter: number of splits:4
22/03/20 08:11:17 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_16
46550420897_0016
22/03/20 08:11:17 INFO impl.YarnClientImpl: Submitted application application_16
46550420897_0016
22/03/20 08:11:17 INFO mapreduce.Job: The url to track the job: http://quickstar
t.cloudera:8088/proxy/application_1646550420897_0016/
22/03/20 08:11:17 INFO mapreduce.Job: Running job: job_1646550420897_0016
22/03/20 08:11:31 INFO mapreduce.Job: Job job_1646550420897_0016 running in uber
mode : false
22/03/20 08:11:31 INFO mapreduce.Job: map 0% reduce 0%
22/03/20 08:12:15 INFO mapreduce.Job: map 25% reduce 0%
22/03/20 08:12:16 INFO mapreduce.Job: map 50% reduce 0%
22/03/20 08:12:18 INFO mapreduce.Job: map 75% reduce 0%
22/03/20 08:12:19 INFO mapreduce.Job: map 100% reduce 0%
22/03/20 08:12:19 INFO mapreduce.Job: Job job_1646550420897_0016 completed succe
ssfully
cloudera@quickstart:~$
```

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$
File Edit View Search Terminal Help
22/03/20 08:12:18 INFO mapreduce.Job: map 75% reduce 0%
22/03/20 08:12:19 INFO mapreduce.Job: map 100% reduce 0%
22/03/20 08:12:19 INFO mapreduce.Job: Job job_1646550420897_0016 completed succe
ssfully
22/03/20 08:12:19 INFO mapreduce.Job: Counters: 31
File System Counters
FILE: Number of bytes read=0
FILE: Number of bytes written=685412
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=481
HDFS: Number of bytes written=60
HDFS: Number of read operations=16
HDFS: Number of large read operations=0
HDFS: Number of write operations=8
Job Counters
Killed map tasks=1
Launched map tasks=4
Other local map tasks=5
Total time spent by all maps in occupied slots (ms)=165135
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=165135
Total vcore-milliseconds taken by all map tasks=165135
Total megabyte-milliseconds taken by all map tasks=169090240
Map-Reduce Framework
Map input records=6
Map output records=6
Input split bytes=481
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=2448
CPU time spent (ms)=6780
Physical memory (bytes) snapshot=500707328
Virtual memory (bytes) snapshot=6040596480
Total committed heap usage (bytes)=243007488
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=0
22/03/20 08:12:19 INFO mapreduce.ImportJobBase: Transferred 60 bytes in 66.7391
seconds (0.899 bytes/sec)
22/03/20 08:12:19 INFO mapreduce.ImportJobBase: Retrieved 6 records.
[cloudera@quickstart ~]$ hadoop fs -ls
Found 2 items
drwxr-xr-x - cloudera cloudera 0 2022-03-20 08:12 department1
drwxr-xr-x - cloudera cloudera 0 2022-03-20 07:59 departments
[cloudera@quickstart ~]$
cloudera@quickstart:~$
```

As we can see 6 records are retrived successfully.

hadoop fs -ls

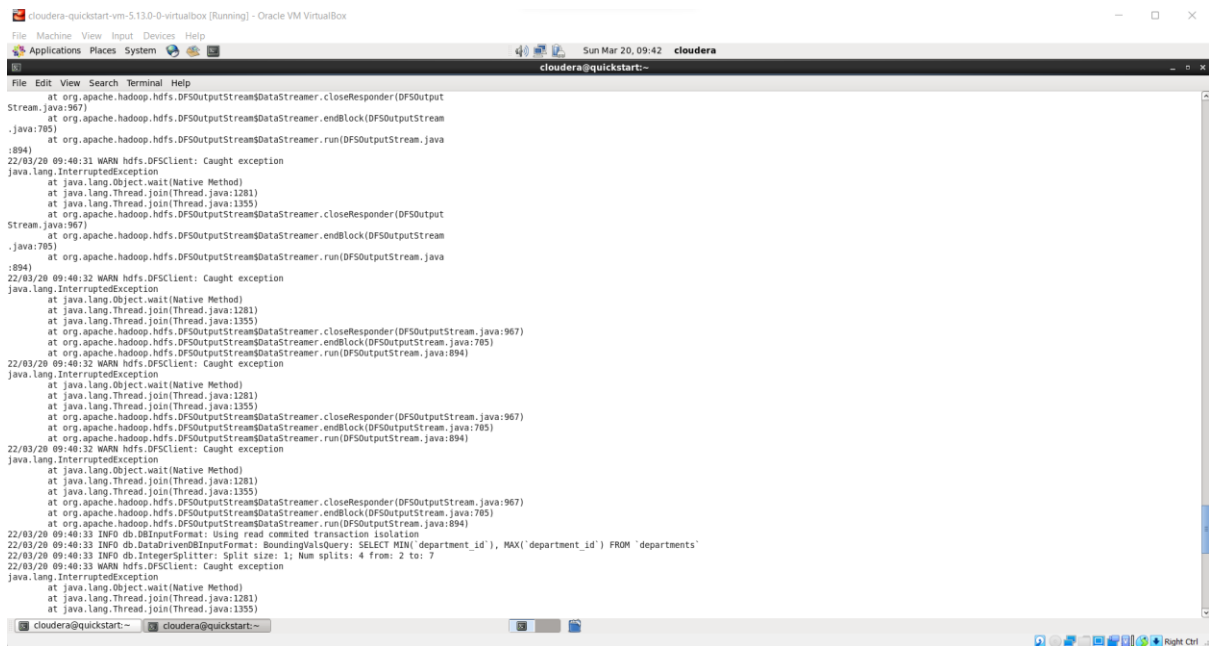
```
[cloudera@quickstart ~]$ hadoop fs -ls
Found 2 items
drwxr-xr-x - cloudera cloudera      0 2022-03-20 08:12 department1
drwxr-xr-x - cloudera cloudera      0 2022-03-20 07:59 departments
[cloudera@quickstart ~]$
```

hadoop fs -cat /user/cloudera/department1/part*

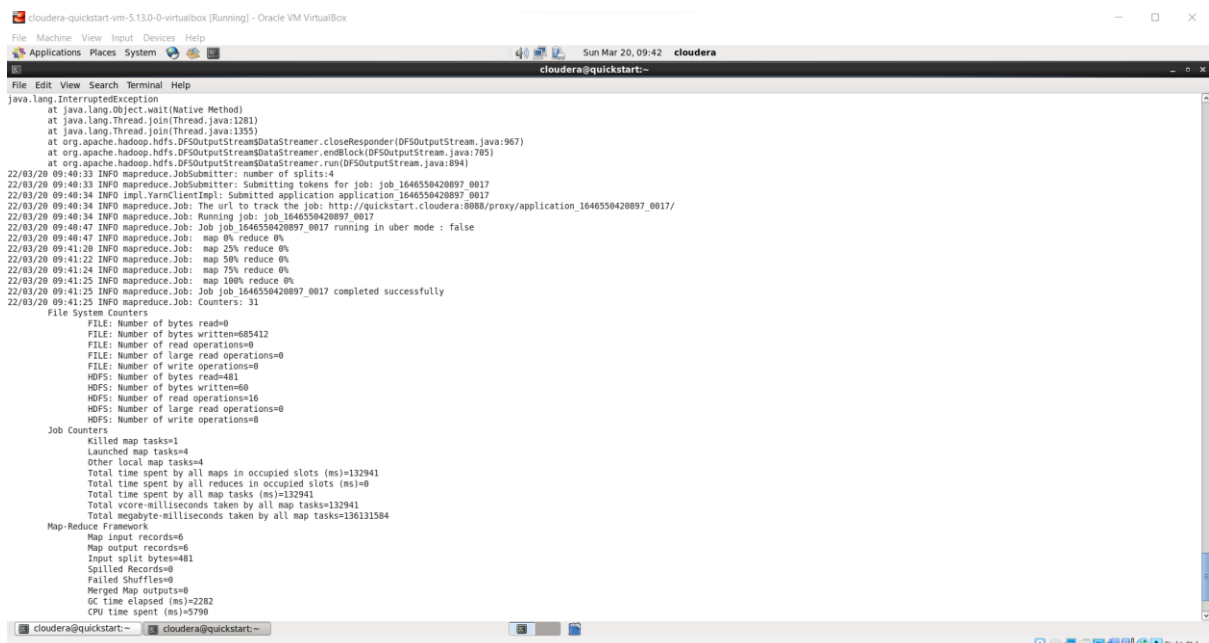
```
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/department1/part*
2,Fitness
3,Footwear
4,Apparel
5,Golf
6,Outdoors
7,Fan Shop
[cloudera@quickstart ~]$
```

sqoop import --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root --table departments --where "department_id>4" --target-dir/user/cloudera/department2;

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
Sun Mar 20, 09:42 cloudera
cloudera@quickstart:~
File Edit View Search Terminal Help
6,Outdoors
7,Fan Shop
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://quickstart:3306/retail_db
--password cloudera --username root --table departments --target-dir /user/cloudera/de
partment2;
Warning: /usr/lib/sqoop/.accumulo does not exist! Accumulo imports will fail.
Please set HADOOP_HOME to the root of your Accumulo installation.
22/03/20 09:40:12 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/03/20 09:40:12 WARN tool.BaseSqoopTool: Setting your password on the command-line is
insecure. Consider using --P instead.
22/03/20 09:40:13 INFO manager.MySQLManager: Preparing to use a MySQL streaming results
et.
22/03/20 09:40:13 INFO tool.CodeGenTool: Beginning code generation
22/03/20 09:40:15 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `de
partments` AS t LIMIT 1
22/03/20 09:40:15 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `de
partments` AS t LIMIT 1
22/03/20 09:40:15 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-ma
preduce
Note: /tmp/sqoop-cloudera/compile/9d4def418e0eab93a333d5208329f5/departments.java use
s or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
22/03/20 09:40:25 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/co
mpile/9d4def418e0eab93a333d5208329f5/departments.jar
22/03/20 09:40:25 WARN manager.MySQLManager: It looks like you are importing from mysql
.
22/03/20 09:40:25 WARN manager.MySQLManager: This transfer can be faster! Use the --dir
ect
22/03/20 09:40:25 WARN manager.MySQLManager: option to exercise a MySQL-specific fast p
ath.
22/03/20 09:40:25 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertT
oNull (mysql)
22/03/20 09:40:25 INFO mapreduce.ImportJobBase: Beginning import of departments
22/03/20 09:40:25 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Ins
tead, use mapreduce.jobtracker.address
22/03/20 09:40:27 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, us
e mapreduce.job.jar
22/03/20 09:40:29 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Inste
ad, use mapreduce.job.maps
22/03/20 09:40:29 INFO Client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/03/20 09:40:31 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DFSOutput
Stream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutputStream
.java:795)
cloudera@quickstart:~
cloudera@quickstart:~
```



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$
File Edit View Search Terminal Help
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DFSOutput
Stream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutputStream
.java:785)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStream.java
:894)
22/03/20 09:40:31 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DFSOutput
Stream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutputStream
.java:785)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStream.java
:894)
22/03/20 09:40:32 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DFSOutput
Stream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutputStream
.java:785)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStream.java
:894)
22/03/20 09:40:33 INFO db.ObjectInputFormat: Using read committed transaction isolation
22/03/20 09:40:33 INFO db.DataDriverInputFormat: BoundingvalQuery: SELECT MIN(' department_id' ), MAX(' department_id' ) FROM ' departments'
22/03/20 09:40:33 INFO db.IntegerSplitter: Split size: 1; Num splits: 4 from: 2 to: 7
22/03/20 09:40:33 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at java.lang.Thread.join(Thread.java:1355)
cloudera@quickstart:~$
```



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~$
File Edit View Search Terminal Help
java.lang.InterruptedException
at java.lang.Object.wait(Native Method)
at java.lang.Thread.join(Thread.java:1281)
at java.lang.Thread.join(Thread.java:1355)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DFSOutput
Stream.java:967)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutputStream
.java:785)
at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStream.java
:894)
22/03/20 09:40:33 INFO mapreduce.JobSubmitter: number of splits:4
22/03/20 09:40:34 INFO Impl.VarnClientImpl: Submitted application application_1646550420897_0017
22/03/20 09:40:34 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8080/proxy/application_1646550420897_0017/
22/03/20 09:40:34 INFO mapreduce.Job: Running job: job_1646550420897_0017
22/03/20 09:40:47 INFO mapreduce.Job: Job job_1646550420897_0017 running in uber mode : false
22/03/20 09:40:47 INFO mapreduce.Job: map 0% reduce 0%
22/03/20 09:41:20 INFO mapreduce.Job: map 25% reduce 0%
22/03/20 09:41:22 INFO mapreduce.Job: map 50% reduce 0%
22/03/20 09:41:24 INFO mapreduce.Job: map 75% reduce 0%
22/03/20 09:41:25 INFO mapreduce.Job: map 100% reduce 0%
22/03/20 09:41:25 INFO mapreduce.Job: Job job_1646550420897_0017 completed successfully
22/03/20 09:41:25 INFO mapreduce.Job: Counters: 31
File System Counters
FILE: Number of bytes read=0
FILE: Number of bytes written=685412
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=481
HDFS: Number of bytes written=60
HDFS: Number of read operations=16
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
Job Counters
Killed map tasks=1
Launched map tasks=4
Other local map tasks=4
Total time spent by all maps in occupied slots (ms)=132941
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=132941
Total vcore-milliseconds taken by all map tasks=132941
Total megabyte-milliseconds taken by all map tasks=136131584
Map-Reduce Framework
Map input records=6
Map output records=6
Input split bytes=481
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=2282
CPU time spent (ms)=5790
cloudera@quickstart:~$
```

```
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=2282
CPU time spent (ms)=5790
Physical memory (bytes) snapshot=483442688
Virtual memory (bytes) snapshot=6040580096
Total committed heap usage (bytes)=243007488

File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=60
22/03/20 09:41:26 INFO mapreduce.ImportJobBase: Transferred 60 bytes in 56.3145 seconds (1.0654 bytes/sec)
22/03/20 09:41:26 INFO mapreduce.ImportJobBase: Retrieved 6 records.
```

As we can see 3 records are retrieved successfully.

hadoop fs -ls /user/cloudera/department2

```
-ls/user/cloudera/department2: unknown command
[cloudera@quickstart ~]$ hadoop fs -ls /user/cloudera/department2
Found 5 items
-rw-r--r--  1 cloudera cloudera      0 2022-03-20 09:41 /user/cloudera/department2
/_SUCCESS
-rw-r--r--  1 cloudera cloudera    21 2022-03-20 09:41 /user/cloudera/department2
/part-m-00000
-rw-r--r--  1 cloudera cloudera    10 2022-03-20 09:41 /user/cloudera/department2
/part-m-00001
-rw-r--r--  1 cloudera cloudera      7 2022-03-20 09:41 /user/cloudera/department2
/part-m-00002
-rw-r--r--  1 cloudera cloudera    22 2022-03-20 09:41 /user/cloudera/department2
/part-m-00003
[cloudera@quickstart ~]$
```

hadoop fs -cat /user/cloudera/department2/part*

```
[cloudera@quickstart ~]$ hadoop fs -cat /user/cloudera/department2/part*
5,Golf
6,Outdoors
7,Fan Shop
[cloudera@quickstart ~]$
```

Now we will see the Export command. So what the export tool does is it will export the data from our hdfs to the RDBMS. So for that we need to have some table in mysql with some records so for that we will now move to mysql.

create table dpt(department_id int not null auto_increment, department_name varchar(50) not null, primary key(department_id));

```
mysql> create table dpt(department_id int not null auto_increment,department_name varchar(50) not null,primary key(department_id));
Query OK, 0 rows affected (0.05 sec)
```

```
mysql>
```

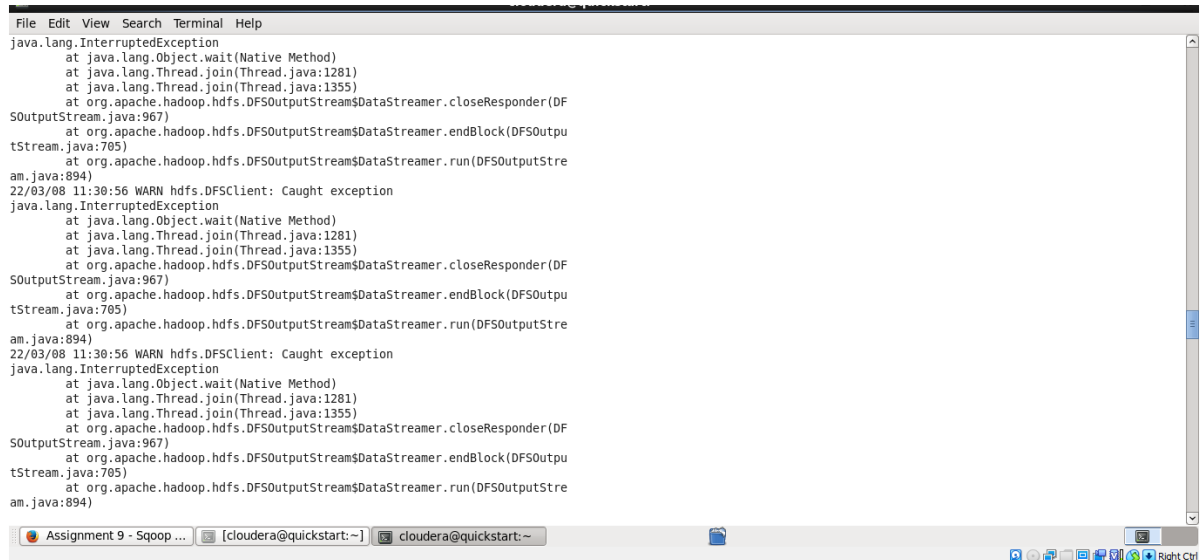
Select * from dpt;

```
mysql> select * from dpt;
Empty set (0.01 sec)

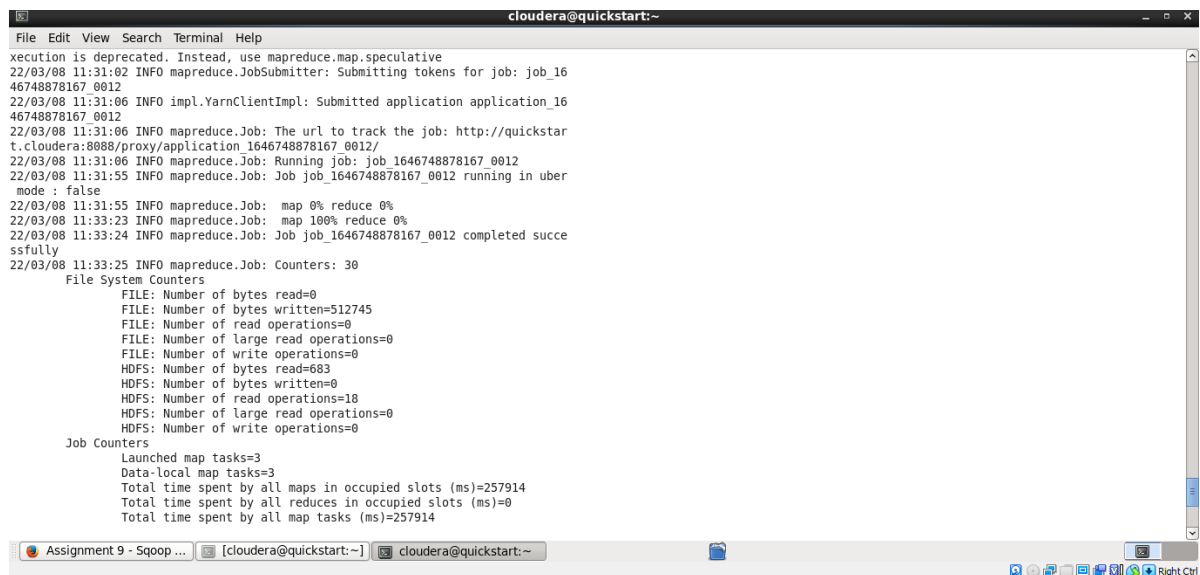
mysql>
```

Now we will exporting the data from the hdfs to dpt table of mysql. Now we will move to the sqoop terminal.

```
sqoop export --connect jdbc:mysql://quickstart:3306/retail_db --password cloudera --username root  
--table dpt --export-dir /user/cloudera/department2;
```



```
File Edit View Search Terminal Help
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DF
SOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutpu
tStream.java:705)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/08 11:30:56 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DF
SOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutpu
tStream.java:705)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
22/03/08 11:30:56 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DF
SOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutpu
tStream.java:705)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStre
am.java:894)
```



```
cloudera@quickstart:~
File Edit View Search Terminal Help
Execution is deprecated. Instead, use mapreduce.map.speculative
22/03/08 11:31:02 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_16
46748878167_0012
22/03/08 11:31:06 INFO impl.YarnClientImpl: Submitted application application_16
46748878167_0012
22/03/08 11:31:06 INFO mapreduce.Job: The url to track the job: http://quickstar
t.cloudera:8088/proxy/application_1646748878167_0012/
22/03/08 11:31:06 INFO mapreduce.Job: Running job: job_1646748878167_0012
22/03/08 11:31:55 INFO mapreduce.Job: Job job_1646748878167_0012 running in uber
mode : false
22/03/08 11:31:55 INFO mapreduce.Job: map 0% reduce 0%
22/03/08 11:33:23 INFO mapreduce.Job: map 100% reduce 0%
22/03/08 11:33:24 INFO mapreduce.Job: Job job_1646748878167_0012 completed succe
ssfully
22/03/08 11:33:25 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=512745
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=683
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=18
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Job Counters
    Launched map tasks=3
    Data-local map tasks=3
    Total time spent by all maps in occupied slots (ms)=257914
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=257914
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
HDFS: Number of bytes written=0  
HDFS: Number of read operations=18  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=0  
Job Counters  
  Launched map tasks=3  
  Data-local map tasks=3  
  Total time spent by all maps in occupied slots (ms)=257914  
  Total time spent by all reduces in occupied slots (ms)=0  
  Total time spent by all map tasks (ms)=257914  
  Total vcore-milliseconds taken by all map tasks=257914  
  Total megabyte-milliseconds taken by all map tasks=264103936  
Map-Reduce Framework  
  Map input records=3  
  Map output records=3  
  Input split bytes=627  
  Spilled Records=0  
  Failed Shuffles=0  
  Merged Map outputs=0  
  GC time elapsed (ms)=4743  
  CPU time spent (ms)=8450  
  Physical memory (bytes) snapshot=363651072  
  Virtual memory (bytes) snapshot=4524171264  
  Total committed heap usage (bytes)=182255616  
File Input Format Counters  
  Bytes Read=0  
File Output Format Counters  
  Bytes Written=0  
22/03/08 11:33:25 INFO mapreduce.ExportJobBase: Transferred 683 bytes in 154.475  
4 seconds (4,4214 bytes/sec)  
22/03/08 11:33:25 INFO mapreduce.ExportJobBase: Exported 3 records.  
[cloudera@quickstart ~]$
```

Select * from dpt;

```
mysql> select * from dpt;  
+-----+-----+  
| department_id | department_name |  
+-----+-----+  
|          5 | Golf            |  
|          6 | Outdoors        |  
|          7 | Fan Shop        |  
+-----+-----+  
3 rows in set (0.00 sec)  
  
mysql>
```

As we can see these 3 records which are present in department2 table are successfully exported inside the dpt table of mysql.

