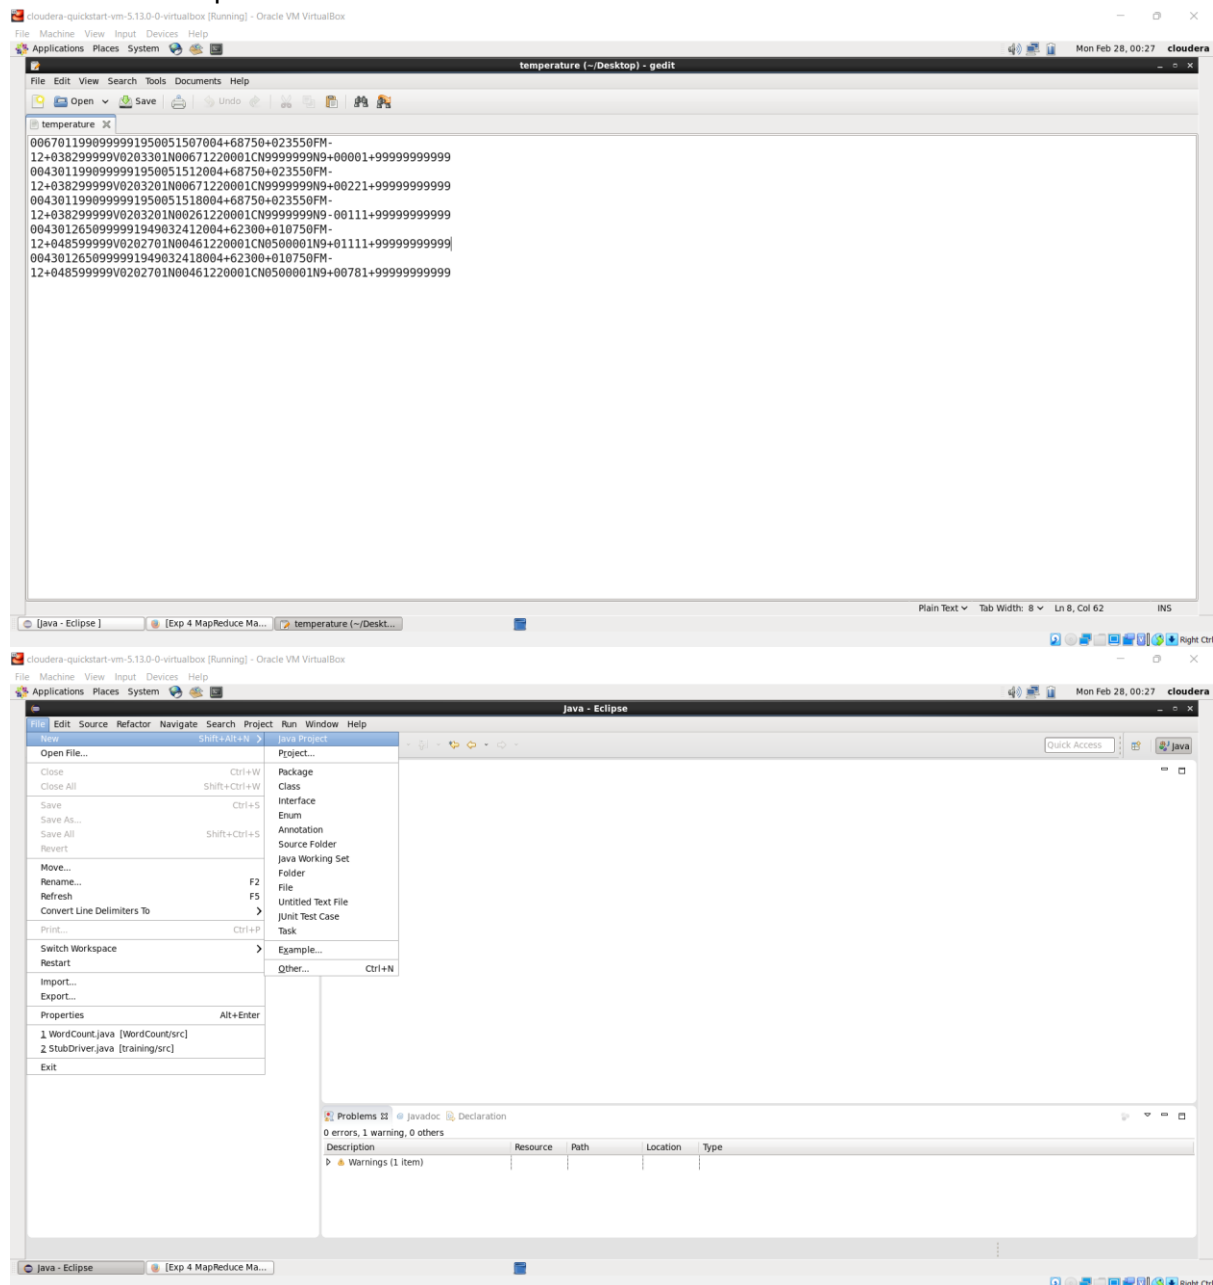
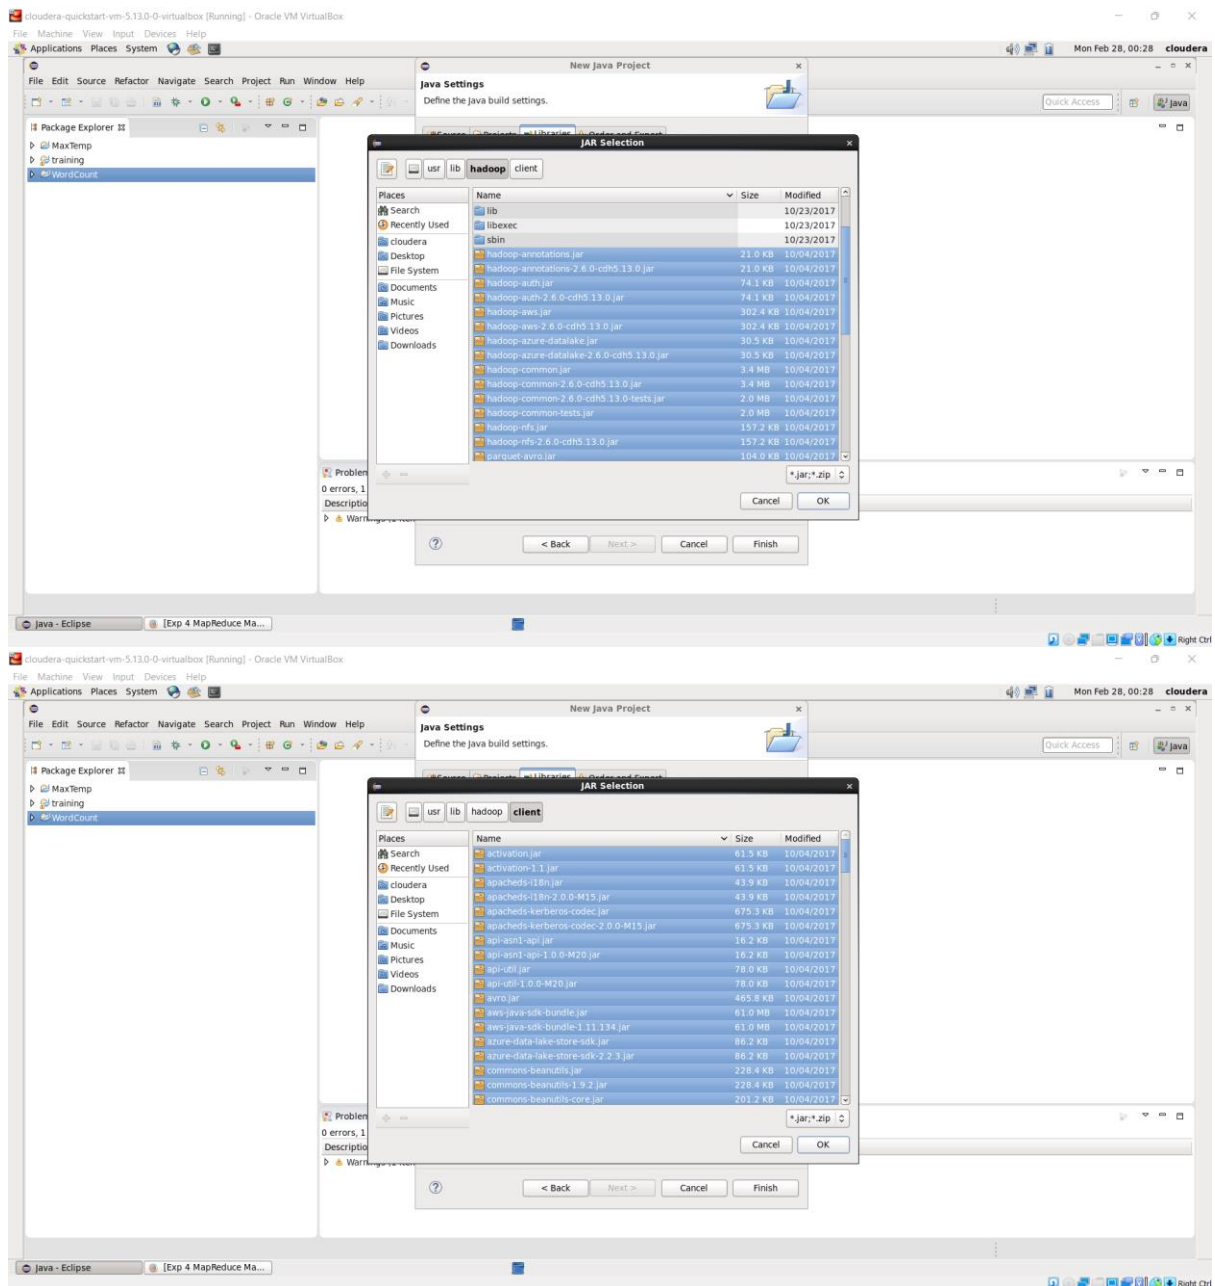


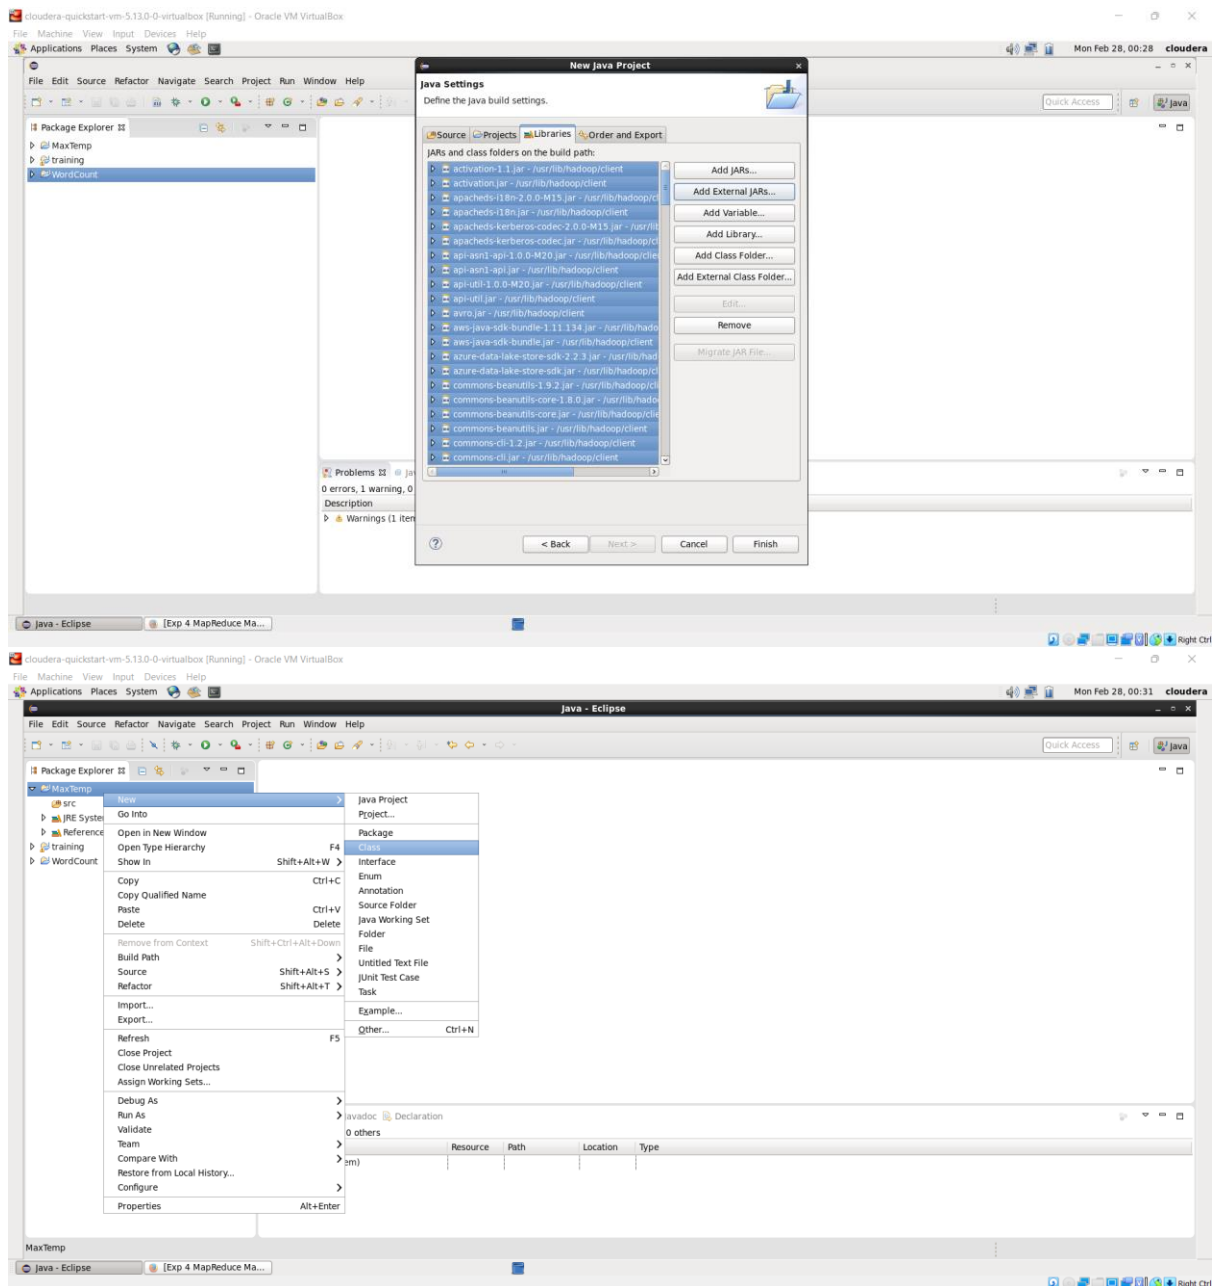
Practical 4

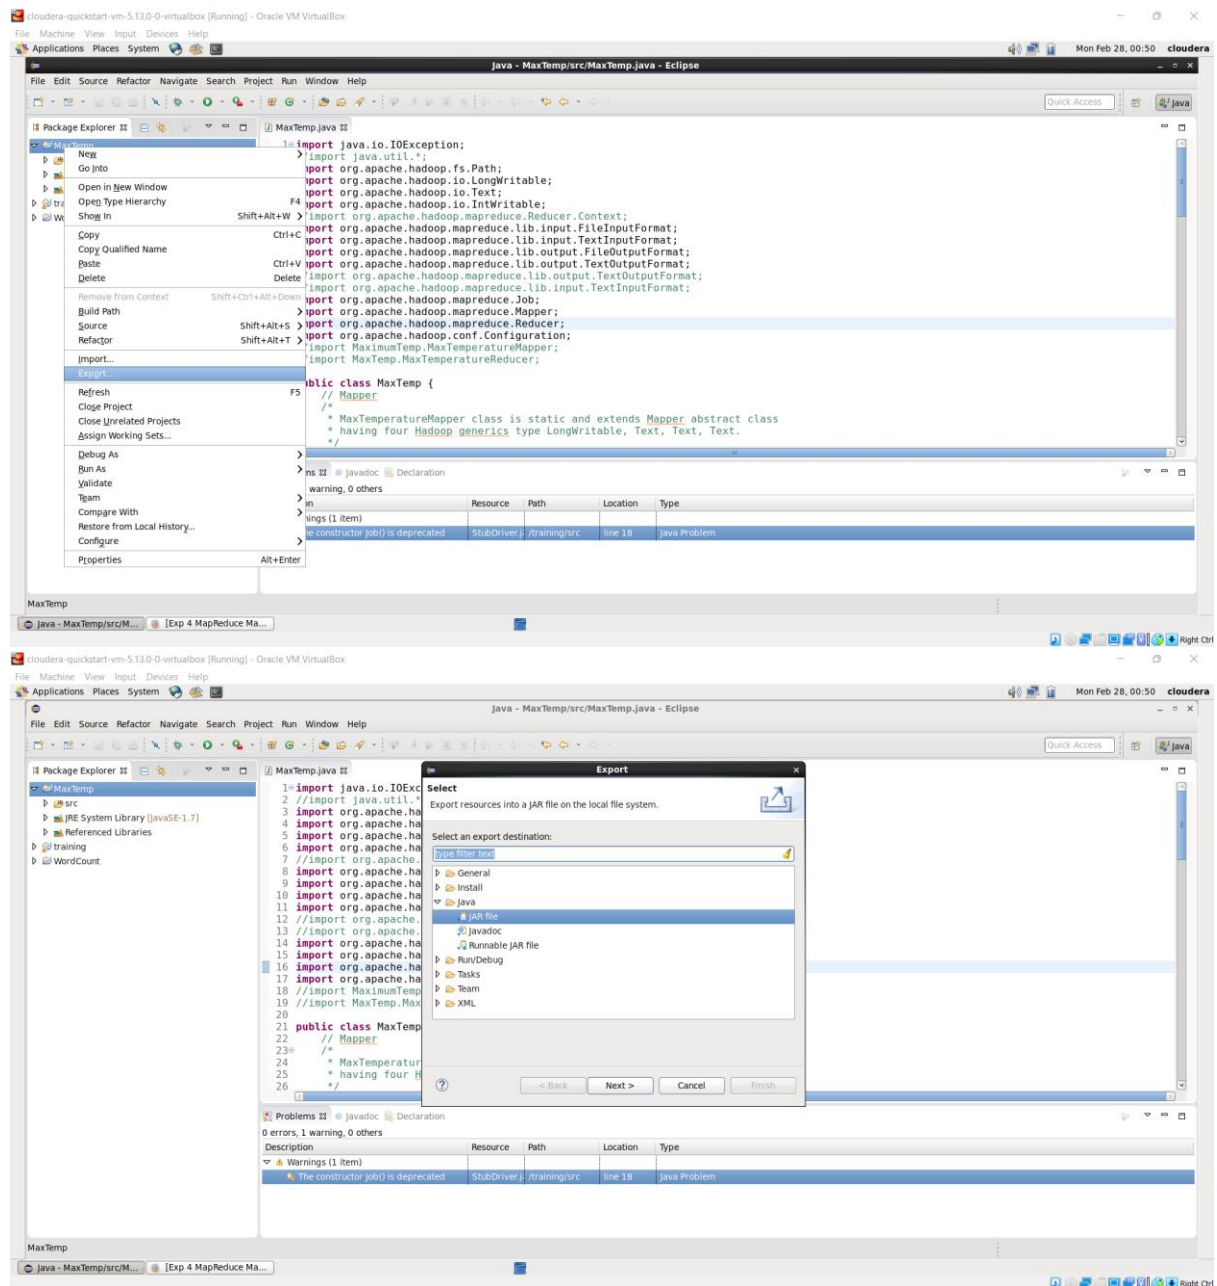
To determine maximum temperature using Hadoop MapReduce

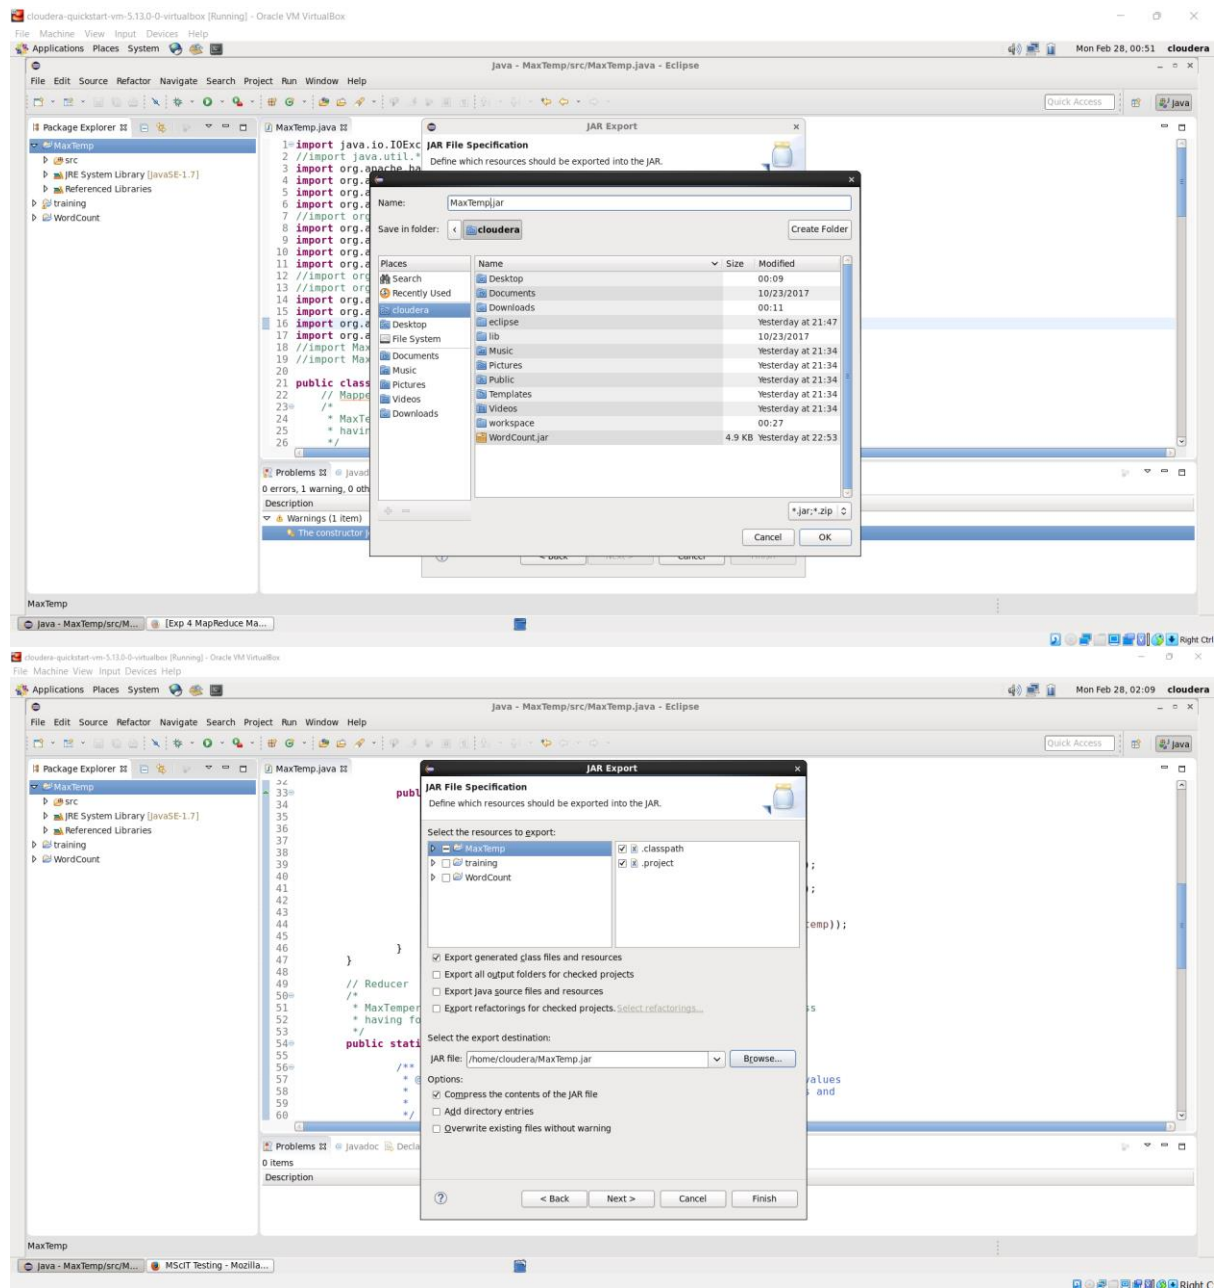
Dataset – temperature.txt







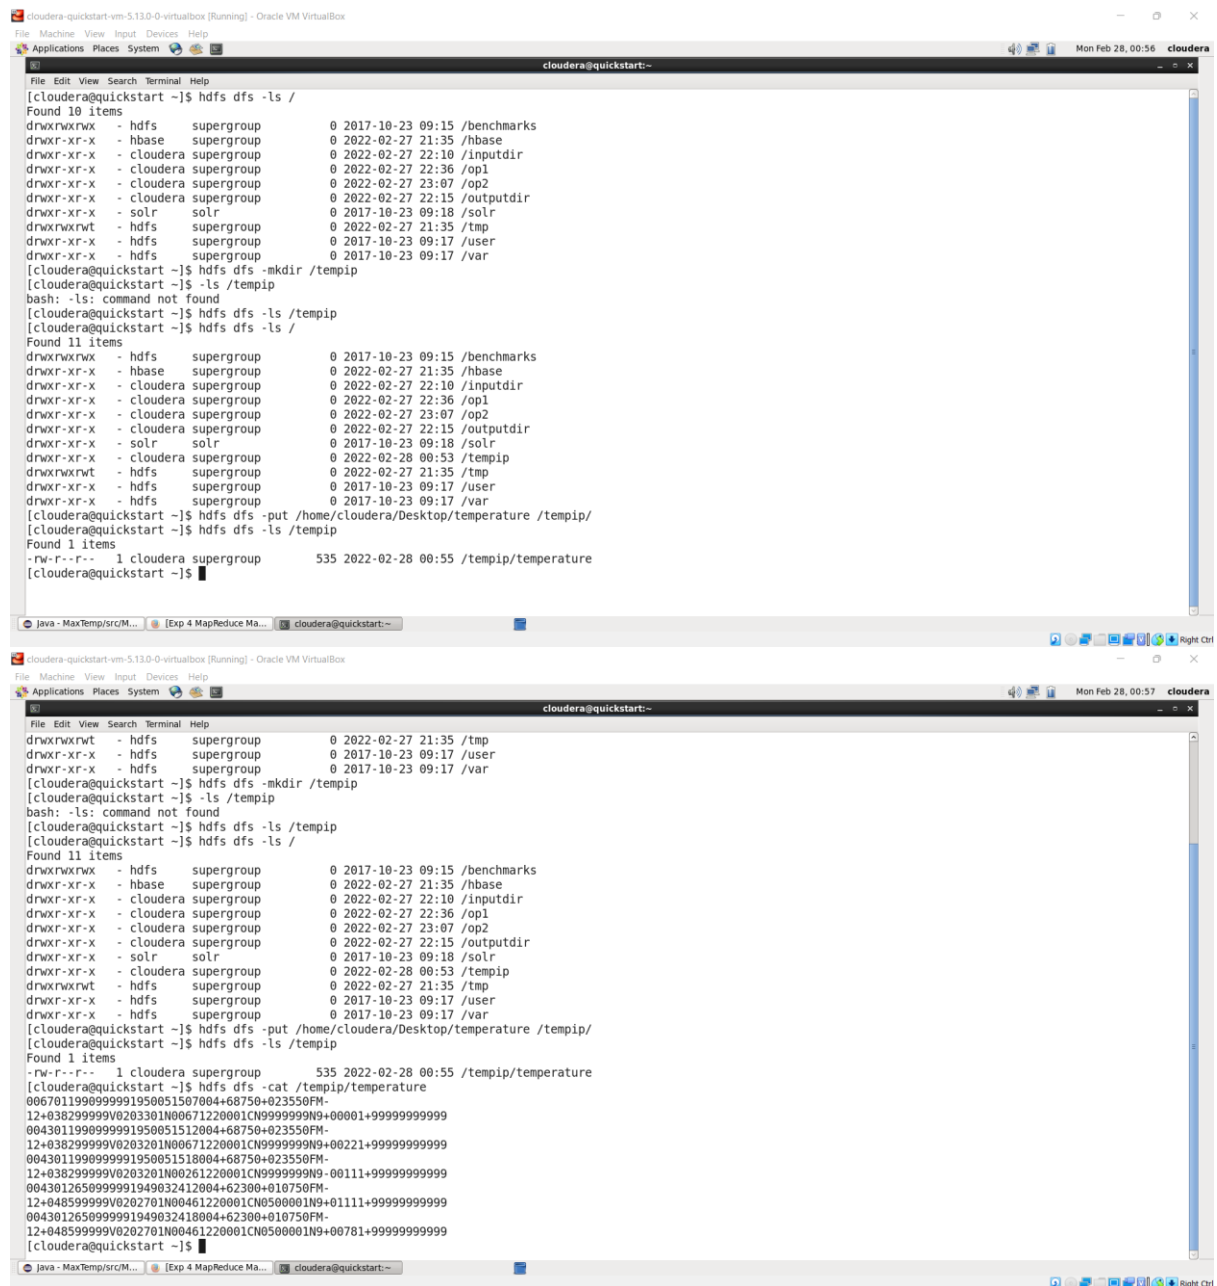




The screenshot shows a terminal window within a Cloudera Quickstart VM. The user runs the command `hdfs dfs -ls /` to list the root directory. The output shows 10 items, including `/benchmarks`, `/hbase`, `/inputdir`, `/op1`, `/op2`, `/outputdir`, `/solr`, `/tmp`, `/user`, and `/var`. The user then runs `hdfs dfs -mkdir /tempip`, which results in a `bash: -ls: command not found` error. The user then runs `hdfs dfs -ls /tempip`, which results in a `bash: -ls: command not found` error. The user then runs `hdfs dfs -ls /` again, which shows 11 items, including `/tempip`. The user then runs `hdfs dfs -ls /tempip`, which shows 11 items, including `/tempip`.

```
cloudera@quickstart:~$ hdfs dfs -ls /
Found 10 items
drwxrwxrwx - hdfs supergroup 0 2017-10-23 09:15 /benchmarks
drwxr-xr-x - hbase supergroup 0 2022-02-27 21:35 /hbase
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:10 /inputdir
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:36 /op1
drwxr-xr-x - cloudera supergroup 0 2022-02-27 23:07 /op2
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:15 /outputdir
drwxr-xr-x - solr solr 0 2017-10-23 09:18 /solr
drwxrwxrwt - hdfs supergroup 0 2022-02-27 21:35 /tmp
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /user
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /var
cloudera@quickstart:~$

cloudera@quickstart:~$ hdfs dfs -mkdir /tempip
bash: -ls: command not found
cloudera@quickstart:~$ hdfs dfs -ls /tempip
bash: -ls: command not found
cloudera@quickstart:~$ hdfs dfs -ls /
Found 11 items
drwxrwxrwx - hdfs supergroup 0 2017-10-23 09:15 /benchmarks
drwxr-xr-x - hbase supergroup 0 2022-02-27 21:35 /hbase
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:10 /inputdir
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:36 /op1
drwxr-xr-x - cloudera supergroup 0 2022-02-27 23:07 /op2
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:15 /outputdir
drwxr-xr-x - solr solr 0 2017-10-23 09:18 /solr
drwxr-xr-x - cloudera supergroup 0 2022-02-28 00:53 /tempip
drwxrwxrwt - hdfs supergroup 0 2022-02-27 21:35 /tmp
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /user
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /var
cloudera@quickstart:~$
```



The image displays two screenshots of a terminal window running on a cloudera-quickstart virtual machine. The terminal shows the execution of HDFS commands to list, create, and write files.

Top Screenshot:

```
cloudera@quickstart:~$ hdfs dfs -ls /
Found 10 items
drwxrwxrwx - hdfs supergroup 0 2017-10-23 09:15 /benchmarks
drwxr-xr-x - hbase supergroup 0 2022-02-27 21:35 /hbase
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:10 /inputdir
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:36 /op1
drwxr-xr-x - cloudera supergroup 0 2022-02-27 23:07 /op2
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:15 /outputdir
drwxr-xr-x - solr solr 0 2017-10-23 09:18 /solr
drwxrwxrwt - hdfs supergroup 0 2022-02-27 21:35 /tmp
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /user
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /var
cloudera@quickstart:~$ hdfs dfs -mkdir /tempip
cloudera@quickstart:~$ -ls /tempip
bash: -ls: command not found
cloudera@quickstart:~$ hdfs dfs -ls /tempip
cloudera@quickstart:~$ hdfs dfs -ls /
Found 11 items
drwxrwxrwx - hdfs supergroup 0 2017-10-23 09:15 /benchmarks
drwxr-xr-x - hbase supergroup 0 2022-02-27 21:35 /hbase
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:10 /inputdir
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:36 /op1
drwxr-xr-x - cloudera supergroup 0 2022-02-27 23:07 /op2
drwxr-xr-x - cloudera supergroup 0 2022-02-27 22:15 /outputdir
drwxr-xr-x - solr solr 0 2017-10-23 09:18 /solr
drwxr-xr-x - cloudera supergroup 0 2022-02-28 00:53 /tempip
drwxrwxrwt - hdfs supergroup 0 2022-02-27 21:35 /tmp
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /user
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /var
cloudera@quickstart:~$ hdfs dfs -put /home/cloudera/Desktop/temperature /tempip/
cloudera@quickstart:~$ hdfs dfs -ls /tempip
Found 1 items
-rw-r--r-- 1 cloudera supergroup 535 2022-02-28 00:55 /tempip/temperature
cloudera@quickstart:~$
```

Bottom Screenshot:

```
cloudera@quickstart:~$ hdfs dfs -ls /tempip
Found 1 items
-rw-r--r-- 1 cloudera supergroup 535 2022-02-28 00:55 /tempip/temperature
cloudera@quickstart:~$ hdfs dfs -cat /tempip/temperature
0067011980999991950051507004+68750+023550FM-
12+038299999V0203301N00671220001CN9999999N9+00001+9999999999
0043011980999991950051512004+68750+023550FM-
12+038299999V0203201N00671220001CN9999999N9+00221+9999999999
0043011980999991950051518004+68750+023550FM-
12+038299999V0203201N00671220001CN9999999N9-00111+9999999999
0043012650999991949032412004+62300+010750FM-
12+048599999V0202701N00461220001CN0500001N9+01111+9999999999
0043012650999991949032418004+62300+010750FM-
12+048599999V0202701N00461220001CN0500001N9+00781+9999999999
cloudera@quickstart:~$
```



```

cloudera@quickstart:~$ hadoop jar /home/cloudera/MaxTemp.jar MaxTemp /temp/temperature /tempop1
22/02/28 02:11:36 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/02/28 02:11:36 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/02/28 02:11:36 INFO input.FileInputFormat: Total input paths to process : 1
22/02/28 02:11:37 INFO mapreduce.JobSubmitter: number of splits:1
22/02/28 02:11:37 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1646026409740_0004
22/02/28 02:11:37 INFO impl.YarnClientImpl: Submitted application application_1646026409740_0004
22/02/28 02:11:37 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1646026409740_0004/
22/02/28 02:11:37 INFO mapreduce.Job: Running job: job_1646026409740_0004
22/02/28 02:11:45 INFO mapreduce.Job: Job job_1646026409740_0004 running in uber mode : false
22/02/28 02:11:45 INFO mapreduce.Job: map 0% reduce 0%
22/02/28 02:11:51 INFO mapreduce.Job: Task Id : attempt_1646026409740_0004_m_000000_0, Status : FAILED
Error: java.lang.StringIndexOutOfBoundsException: String index out of range: 87
    at java.lang.String.charAt(String.java:658)
    at MaxTemp$MaxTemperatureMapper.map(MaxTemp.java:38)
    at MaxTemp$MaxTemperatureMapper.map(MaxTemp.java:1)
    at org.apache.hadoop.mapreduce.Mapper.run(Mapper.java:145)
    at org.apache.hadoop.mapred.MapTask.runNewMapper(MapTask.java:793)
    at org.apache.hadoop.mapred.MapTask.run(MapTask.java:341)
    at org.apache.hadoop.mapred.YarnChild$2.run(YarnChild.java:164)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:415)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1917)
    at org.apache.hadoop.mapred.YarnChild.main(YarnChild.java:158)

22/02/28 02:11:55 INFO mapreduce.Job: Task Id : attempt_1646026409740_0004_m_000000_1, Status : FAILED
Error: java.lang.StringIndexOutOfBoundsException: String index out of range: 87
    at java.lang.String.charAt(String.java:658)
    at MaxTemp$MaxTemperatureMapper.map(MaxTemp.java:38)
    at MaxTemp$MaxTemperatureMapper.map(MaxTemp.java:1)
    at org.apache.hadoop.mapreduce.Mapper.run(Mapper.java:145)
    at org.apache.hadoop.mapred.MapTask.runNewMapper(MapTask.java:793)
    at org.apache.hadoop.mapred.MapTask.run(MapTask.java:341)
    at org.apache.hadoop.mapred.YarnChild$2.run(YarnChild.java:164)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:415)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1917)
    at org.apache.hadoop.mapred.YarnChild.main(YarnChild.java:158)

22/02/28 02:12:00 INFO mapreduce.Job: Task Id : attempt_1646026409740_0004_m_000000_2, Status : FAILED
Error: java.lang.StringIndexOutOfBoundsException: String index out of range: 87
    at java.lang.String.charAt(String.java:658)
    at MaxTemp$MaxTemperatureMapper.map(MaxTemp.java:38)
    at MaxTemp$MaxTemperatureMapper.map(MaxTemp.java:1)
    at org.apache.hadoop.mapreduce.Mapper.run(Mapper.java:145)
    at org.apache.hadoop.mapred.MapTask.runNewMapper(MapTask.java:793)
    at org.apache.hadoop.mapred.MapTask.run(MapTask.java:341)
    at org.apache.hadoop.mapred.YarnChild$2.run(YarnChild.java:164)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:415)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1917)
    at org.apache.hadoop.mapred.YarnChild.main(YarnChild.java:158)

22/02/28 02:12:07 INFO mapreduce.Job: map 100% reduce 100%
22/02/28 02:12:07 INFO mapreduce.Job: Job job_1646026409740_0004 failed with state FAILED due to: Task failed task_1646026409740_0004_m_000000
Job failed as tasks failed. failedMaps:1 failedReduces:0

22/02/28 02:12:07 INFO mapreduce.Job: Counters: 10
    Job Counters
      Failed map tasks=4
      Killed reduce tasks=1
      Launched map tasks=4
      Other local map tasks=3
      Data-local map tasks=1
      Total time spent by all maps in occupied slots (ms)=13732
      Total time spent by all reduces in occupied slots (ms)=0
      Total time spent by all map tasks (ms)=13732
      Total vcore-milliseconds taken by all map tasks=13732
      Total megabyte-milliseconds taken by all map tasks=14061568
cloudera@quickstart:~$

```